

# STF Specification

Simple Trace Format Specification

Version 1.3

## MIT License

Copyright (c) 2020 sparcians

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

# Contents

[Contents](#)

[ChangeLog](#)

[Introduction](#)

[Use Cases](#)

[What is Standardized](#)

[What is NOT Standardized](#)

[How to read this specification](#)

[Organization of the Specification](#)

[STF File Anatomy](#)

[STF Record Specification](#)

[Record Format](#)

[Specification of Records](#)

[STF Exception Event and Event Meta-Data](#)

[Terminology / Definitions](#)

# ChangeLog

This section captures major incremental changes, and what motivated the changes.

| Motivation/Intent  | What changed   |
|--|--|
| <b>Version 1.3</b>   |  |
| Open Sourcing Prep   | Cosmetic changes, rewrite of the introductory paragraph, and rewrite of introduction to STF file format section.   |
| <b>Version 1.2</b>   |  |
| InstExceptionPCTarget record covers a subset of what EventRecord and the associated EventPCTargetRecord cover per STF Exception Table definitions.   | Removed InstExceptionPCTargetRecord<br><br>Breaks backward compatibility   |
| Provide a mechanism to record transitions between execution privilege modes (User, Supervisor, Hypervisor, Machine)  | Added STF_EVENT_MODE_CHANGE as a meta-event for capturing transitions into a new execution mode.<br>Modified EventPCTargetRecord to be optional, because this meta-event does not have a PC change associated with it.   |
| Reduce specification complexity and improve processing efficiency by combining two Event records which always go together.   | Merged event meta-data fields from EventContentRecord into the EventRecord, and removed EventContentRecord.<br><br>Breaks backward compatibility   |
| <b>Version 1.1</b>   |  |
| Provide a summary of changes to make it easier for developers to update tools and comply with the new spec.  | Added ChangeLog section  |
| There is a benefit, purely from implementation standpoint, to have Records which precede other Records have lower numerical values. This allows for ordered traversal through related records, ending with the highest number record, concluding a group of records. This is most frequently done with Instruction Encoding (a.k.a. opcode) Records.<br>Recognising that we can easily get into a bind | Records which have “precede” in their semantic definition are set to have lower numbers than the records which follow them, and conversely records which have “follow” in their semantic definition are set to have higher numbers than the records which come before them.<br><br>Breaks backward compatibility |

|  |  |
|--|--|
| with numbering, we moved the Instruction Encoding records to a much higher enum, and moved PTE ahead of the memory accesses. |  |
| PTE records provide translation, so they should come before the memory accesses records.                                     | <p>PTE records shall precede memory accesses Records.</p> <p>Breaks backward compatibility</p>   |
| Shuffled ID's for ordering which seems to be more logical.   | <p>Moved change of flow records to be at the top as we want them to appear first, followed by register changes, and then address translation (table walks), memory accesses, external events, microops, and finally instruction encodings (opcodes)</p> <p>Breaks backward compatibility</p>   |
| Events can have varying amounts of data that need to be expressed.   | <p>EventContentRecord now allows for an extensible number of fields. Definitions of the fields are provided in the STF Exception Event and Event Meta-Data table.</p> <p>It is placed after the EventRecord, to allow potentially adding another related record in the future if need arises.</p> <p>Breaks backward compatibility</p> |
| Support capturing vector register contents   | Added vector type for registers, so that contents can be captured. This is the first step in vector instructions support.  |
| <b>Version 1.0</b>   |  |
| Initial Spec   | Initial Spec   |

# Introduction

Simple Trace Format (STF) is a binary file format for storing instruction traces, agnostic of instruction set architecture. The format defines a standard for capturing information related to instructions, associated register values, memory access addresses and data associated with them, as well as additional context information such as page table walk, interrupts, bus/fabric transaction addresses/data, etc. This document also specifies how tools generating and/or modifying traces can express information which can aid the tools consuming the traces to interpret them appropriately.

## Use Cases

STF is generated by producer tools, such as functional models or hardware, and consumed by tools such as trace-driven performance models, trace analysis tools, and hardware (if the trace contains the relevant functional information).

## What is Standardized

The specification covers the following aspects:

- Header that provides enough context to allow proper interpretation of the remainder of the trace
- Entities that describe the execution environment for the program (instruction encoding mode, privilege level, process ID, etc.)
- Entities that describe the per-instruction attributes and changes to (non-memory) program state (instruction encoding, source register values, destination register values, side-effect changes, synchronous exceptions)
- Entities that describe the per-instruction attributes and changes to memory program state
- Entities that capture non-instruction-based changes to program state (exceptions, external interrupts, non-traced TLB changes, etc.)
- Data syntax of each entity including their bit-field encoding
- Data semantics of each entity
- Relationships among entities

## What is NOT Standardized

The specification does NOT cover the following aspects:

- Instruction Set Disassembly - external binutils helper packages are used to disassemble instructions at run-time

- Compression format - file compression format is chosen independently of this specification from a plethora of existing compression formats and accompanying tools. The standard, however, requires that the API be extensible to allow developers to use his/her own compression formats.

## How to read this specification

- "SHALL" clause - mandatory
- "MUST" clause - mandatory
- "MAY" clause - optional
- "width" - refers to size of data in bits
- *Instruction encoding vs Opcode*:
  - *Instruction encoding* refers to the entire instruction encoding (e.g. 32 bits)
  - *Opcode* refers to the 7 least significant bits [6:0] of the *Instruction Encoding* (per RV32I)
- *Instruction record* is synonymous to *instruction encoding record* (instruction is identified by its instruction encoding record)

## Organization of the Specification

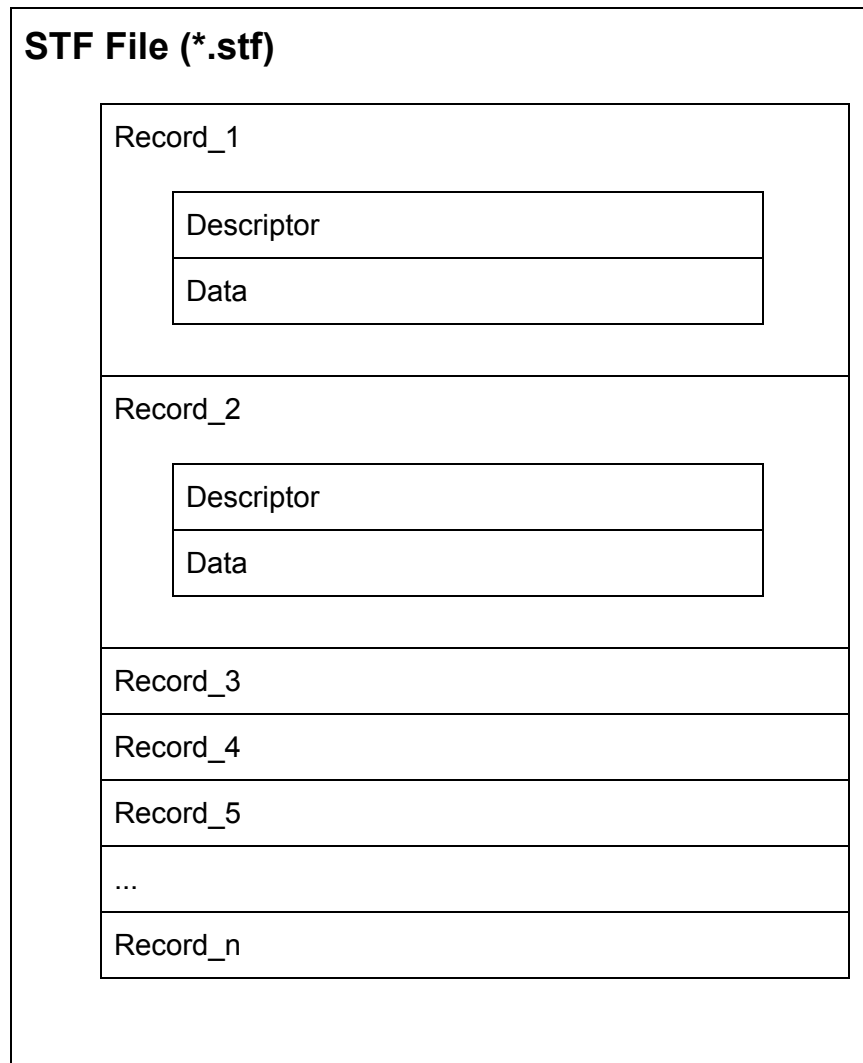
STF specification has 2 domains:

- Syntactic – This part of the specification addresses syntax of entities in an STF file - types of data and associated data structures. It does NOT address the correctness of the relationships among data entities
- Semantic – This part of the specification addresses two main areas:
  - Understanding of each data entity and clarify any ambiguity in interpretation. This area often has to do with the state or pre-condition of the data entities being injected
  - The inter-relationships among various entities of data

# STF File Anatomy

STF **record** is an atomic container which holds a fixed width **descriptor** and a variable width **data** associated with the record. Though the data has variable width, it's structure is well defined by the specification.

An STF file is composed of STF records. Basic anatomy of an STF file is shown in this figure:



Record is always a complete atomic entity.

An STF file shall have a minimum of one record.



Multiple records can be attributed to an instruction. Attributing a set of records to an instruction follows general principles:

- All records, except an instruction encoding record (STF\_INST\_16 or STF\_INST\_32), that precede an instruction encoding record are attributed to the instruction
- There are exceptions to the above rule when it comes to event related records (STF\_EVENT, STF\_EVENT\_CONTENT). Events due to various (mostly asynchronous) exceptions may not be triggered by the instruction they are attributed to according to the STF file.

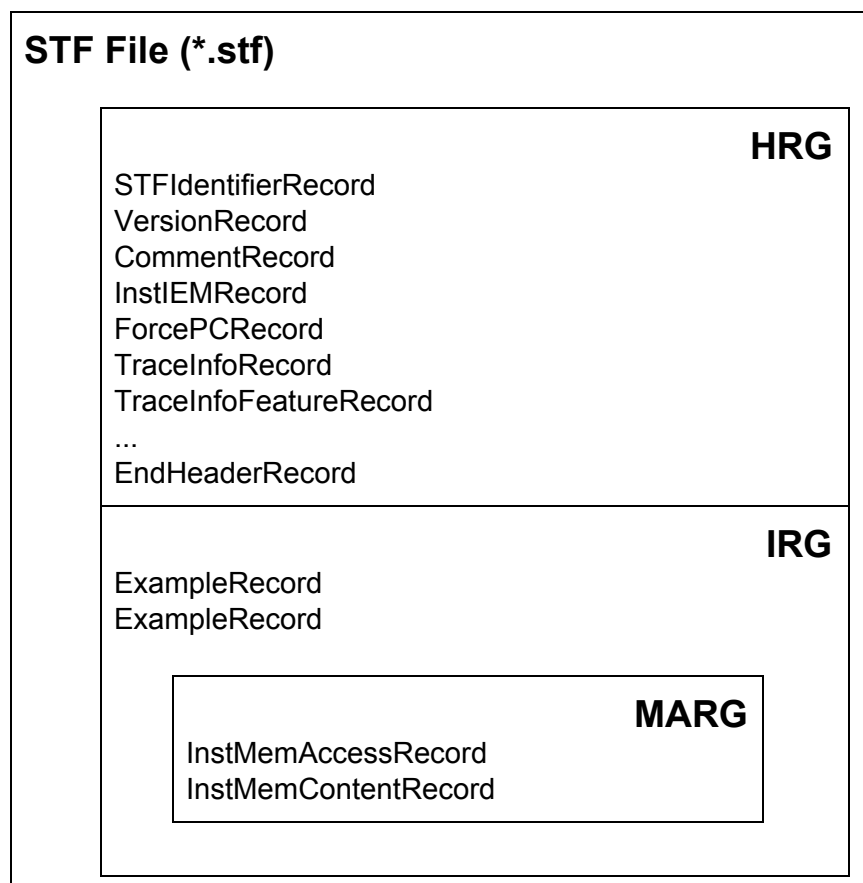
*Header Record Group (HRG)* - Group of record providing information at a trace level.

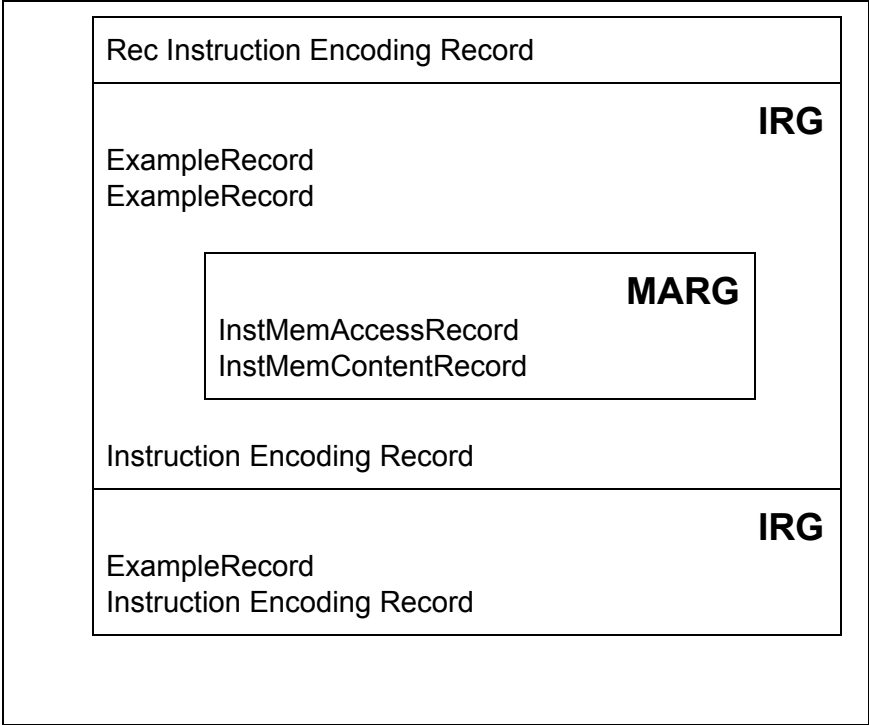
EndHeaderRecord completes and HRG. There shall be only one HRG at the start of a file. If an STF file is modified by a tool, this tool shall update all relevant fields of the HRG (e.g. stitching).

*Instruction Record Group (IRG)* - Group of records attributed to an instruction (i.e. all of the records after previous and before the *instruction encoding record* which completes it's IRG)

*Memory Access Record Group (MARG)* - Group of records associated with a memory access, at minimum including Memory Address and Content.

Relationship between IRG, MARG, various Records, and Instruction Encoding Record is shown in the following figure:





# STF Record Specification

The following table specifies the encoding and semantics of valid STF records.

Example on how to read the table for the `CommentRecord`:

- Record Data Structure: type name is *CommentRecord*, included in the API header `stf.h`
- Record Descriptor: *STF\_COMMENT* descriptor string from the API header `stf.h`. The number in parenthesis () is the enumerated sequence number of the descriptor.
- Mandatory: The field is used for specifying if the record is mandatory. If a record is not mandatory in all cases, this field describes which condition triggers the record. It is mandatory to have a *CommentRecord*.
- Record Data Encoding: Data fields are in series, where the first number within brackets [] is the data field size in bytes. *CommentRecord* encodes the number of bytes of comment data = n bytes into the first 32-bit field, followed by the n-byte long actual string data field.
- Semantics & Additional Comments: Semantics of the specification, as well as free form comments providing details, use cases, and clarification for the record specification.

## Record Format

|   |  |
|---|--|
| <b>Record Data Structure Type</b> see API header <code>stf.h</code><br><b>STF RECORD DESCRIPTOR ()</b> 1-byte wide descriptor with () enumerated value - may be discontinuous |  |
| <b>Mandatory</b><br><br>Specifies if the record is mandatory, and under what conditions.  | <b>Record Data Encoding</b><br>b[bit-span] bit encoding,<br>b[-1] record has only a descriptor with no Record Data (zero bit-encoding) |
|   | <b>Semantics &amp; Additional Comments</b><br>Semantic specification, use-cases, accompanying information, etc.                        |

# Specification of Records

|                                |  |
|--------------------------------|--|
| <b>N/A</b><br>STF_RESERVED (0) |  |
|                                | b[-1] Record has no data, just the descriptor. |
|                                | Reserved for error detection                   |

|  |   |
|--|---|
| <b>STFIdentifierRecord</b><br>STF_IDENTIFIER (1) |   |
| YES  | b[23:0] = "STF" expressed as the magic number   |
|  | <p>Shall exist as the first record of the file.</p> <p>Every tool that creates/modifies stf file shall ensure STFIdentifierRecord exists and is the first record of the file.</p> <p>Display magic number spelling out "STF" at the beginning of the file, to enable identification of STF files without STF tools.</p> |

|   |  |
|---|--|
| <b>VersionRecord</b><br>STF_VERSION (2) |  |
| YES                                     | b[31:0] Major version number of the STF specification<br>b[63:32] Minor version number of the STF specification  |
|   | <p>Shall exist in all STF files as the second record.</p> <p>Every tool that creates/modifies an stf file shall append/update the VersionRecord indicating compliance to a particular STF specification version.</p> <p>Version numbers are maintained in API header stf.h</p> |

|   |  |
|---|--|
| <b>CommentRecord</b><br>STF_COMMENT (3) |  |
| YES                                     | b[31:0] Size of comment string = n bytes<br>b[(8*n+31):32] Comment string data   |
|   | <p>Comment string data is non-null terminated.</p> <p>Every tool that creates/modifies stf file shall append a CommentRecord</p> <p>Shall include git SHAs of all projects used to build the trace generator, and names and versions of converter tools tools for reproducibility.</p> <p>CommentRecord may be added at any point in the trace.</p> <p>Use-case:</p> <p>Log git information about projects used to produce the tracing environment</p> |

|  |   |
|--|---|
|  | Trace converter tools (e.g. trim, morph) - append name and version of STF converter tool that transformed the trace since original raw trace generation |
|--|---|

| <b>ISARecord</b><br>STF_ISA (4) |   |
|---------------------------------|---|
| YES                             | b[15:0] ISA<br>0 = Reserved<br>1 = RISC-V<br>2 = ARM<br>3 = x86<br>4 = Power  |
|                                 | Every tool that creates/modifies stf file shall append/update the ISARecord<br>Shall precede InstIEMRecord, to accomodate provisioning for different Instruction Encoding Modes (IEM) for different ISAs. |

| <b>InstIEMRecord</b><br>STF_INST_IEM (5) |   |
|--|---|
| YES                                      | b[15:0] Instruction encoding mode<br>0 = Reserved<br>1 = STF_INST_IEM_RV32<br>2 = STF_INST_IEM_RV64   |
|  | Every tool that creates/modifies stf file shall append/update theInstIEMRecord.<br>Shall precede the very first instruction encoding record and every instruction encoding change (STF_INST_16 or STF_INST_32) Note: other instruction encoding lengths (e.g.48-bit are not supported at this time))<br>Instruction encoding mode interpretation will depend on the ISA being traced, as captured in the ISARecord. |

| <b>TraceInfoRecord</b><br>STF_TRACE_INFO (6) |  |
|--|--|
| YES  | Information about a trace generator or trace modifier:<br>b[7:0] Trace generator/modifier name:<br>e.g. spike, imperas, sail<br>b[15:8] = major version<br>b[23:16] = minor version<br>b[31:24] = minor minor version<br>b[47:32] Size of comment string = n bytes<br>b[(8*n+48):48] Comment string data |
|  | Every tool that creates/modifies stf file shall append a TraceInfoRecord.  |

## TraceInfoFeatureRecord

STF\_TRACE\_INFO\_FEATURE(7)

|     |  |
|-----|--|
| YES | b[63:0] Features Supported   |
|     | Every tool that creates/modifies stf file shall append a TraceInfoFeatureRecord.<br>Provides information for all features supported or not supported in this trace.<br>When traces are manipulated by tools there should be consistency checking of supported features (e.g. trace stitching). |

## ProcessIDExtRecord

STF\_PROCESS\_ID\_EXT (8)

|     |   |
|-----|---|
| YES | b[31:0] TGID/PID<br>b[63:32] TID<br>b[95:64] ASID   |
|     | Shall precede any instruction record that represents a change from prior instruction record in any of the listed ID's<br>Every tool that creates/modifies stf file shall append/update the ProcessIDExtRecord<br>Indicates any change in thread, process, or address space ID |

## ForcePCRecord

STF\_FORCE\_PC (9)

|     |   |
|-----|---|
| YES | b[63:0] = Virtual address of PC when program COF happens due to non-deterministic cases   |
|     | Shall precede the very first instruction encoding record, and be emitted for any program change of flow (COF) due to non-deterministic cases.<br>Every tool that creates/modifies stf file shall append a ForcePCRecord to indicate the starting PC address.<br>This record indicates virtual address of PC when program change of flow (COF) happens due to non-deterministic cases<br>The next STF_INST_16 / STF_INST_32 record's virtual address will match this record's virtual address<br>Trace tools shall output the architectural PC (even if the low/high bits are non-zero - e.g. ARM).<br>Some architectures may add behaviours based on the low/high bits.<br>Note this record is always 64 bit, even when running in 32bit IEM.<br>stf_dump shows this record as a standalone line item prefix "FORCE_PC" |

## EndHeaderRecord

STF\_END\_HEADER (19)

|     |  |
|-----|--|
| YES | b[-1] Record has no data, just the descriptor.                       |
|     | Every tool that creates an stf file shall append the EndHeaderRecord |

|  |  |
|--|--|
|  | Shall be issued as the last header record in an STF file. Its only purpose is to complete the Header Record Group. |
|--|--|

## InstPCTargetRecord

STF\_INST\_PC\_TARGET (31)

|   |  |
|---|--|
| Whenever COF condition exists due to a branch | b[63:0] Virtual address of target PC when current instruction's branch is taken, causing a change-of-flow (COF).   |
|   | Shall be emitted only for branch based COF (not an exception based COF).<br>Omit this record if the branch is not taken.<br>stf_dump shows this records with a prefix of "PC " |

## InstRegRecord

STF\_INST\_REG (40)

|     |  |
|-----|--|
| YES | b[15:0] Register number. Encoding of the register is outlined in stf_reg_def.h as an enum of type STF_REG. See the register encoding table.<br>b[23:16] Register type encoding:<br>b[19:16] Register type:<br>0000 = reserved<br>0001 = integer<br>0010 = floating point<br>0011 = vector<br>0100 = CSR<br>b[21:20]: Register operand type:<br>00 = reserved<br>01 = state<br>10 = source register<br>11 = destination register<br>b[23:22] reserved<br>b[87:24] Register value/content  |
|     | Description and content of a register relevant to an instruction.<br>Multiple records are used to convey the state of all/required set of registers for trace consumers.<br>Use-case:<br>Functional model uses this record to dump register state if periodic register dump is set (which is the case for default Functional model run) or specific condition for register dump is met (i.e. at the very beginning of a trace)<br>Periodic register dumps are used as reference points in stf2elf flow to speed up mid-trace machine register state determination. |

## InstReadyRegRecord

STF\_INST\_READY\_REG (41)

|    |   |
|----|---|
| NO | b[15:0] Register number of ready register |
|----|---|

|  |  |
|--|--|
|  | <p>This is used when we artificially modify dependencies between instructions.</p> <p>Mark destination register as ready</p> |
|--|--|

| <b>PageTableWalkRecord</b><br>STF_PAGE_TABLE_WALK (50)                           |   |
|--|---|
| Whenever<br>STF_CONT<br>AIN_PTE<br>bit is set in<br>STF_TRACE<br>_INFO<br>record | <p>b[63:0] Virtual address of the page being accessed</p> <p>b[127:64] Instruction count, count starts at index = 0 based (from the beginning of the trace, that encountered the first memory access occurrence from this page)</p> <p>b[159:128] Page size (Functional model definition: "size of translated page in bytes" )</p> <p>b[167:160] Number of PTEs accessed by table walk</p> <p>b[295:168] PTE 0</p> <p>b[231:168] Physical address of PTE</p> <p>b[295:232] Raw PTE (including page attributes,etc.)</p> <p>...</p> <p>b[128*n+295:128*n+168] PTE n</p>  |
|  | <p>This record captures the page table entries (PTEs) accessed during a page table walk. This is a variable length record dependent on the depth of the page table walk. The last PTE in the record should be the leaf PTE that provides the memory translation.</p> <p>Shall precede any new memory access to a page, where new memory access has one or more of the following meanings:</p> <ul style="list-style-type: none"> <li>o The very first memory access, since the beginning of the trace, to an address (virtual) that belongs to a new page</li> <li>o An access to a memory address (virtual) which has been accessed before, but has had an update to the virtual to physical page mapping since.</li> </ul> <p>If there are multiple pages being accessed by an instruction meeting any of the "new memory access" requirements, then multiple page table walk records corresponding to those accesses shall precede the instruction record.</p> <p>There shall be a page table walk record every time a unique page mapping changes or is newly introduced. Unique page mapping is uniquely identifiable by vmid, asid, and VA. i.e. when vmid and asid are not available, the trace format does not recognize PTE information, so we won't write it.</p> <p>For RV32 only valid modes are Bare (no translation) and Sv32 (32 bit Virtual Addressing)</p> <p>For RV64, besides Bare, Sv39 and Sv48 are supported</p> <p>Attributes [7:0] are mapped the same for all cases, but the other bits are different between S32 vs. S39/S48</p> <p>Attributes indicate, among other things, if this is a leaf PTE or not</p> |

| <b>InstMemAccessRecord</b><br>STF_INST_MEM_ACCESS (60) |   |
|--|---|
| YES  | <p>b[63:0] Virtual address of the target memory being accessed by current instruction (read or write).</p> <p>b[79:64] memory access data size</p> <p>b[95:80] memory access attributes.</p> <p>Access type encoding:</p> <p>ToDo: List all access type encodings here (e.g. non-cacheable, write back write allocate, etc.) as well as encodings for prefetches, preloads, etc.</p> <p>b[103:96] access type</p> |



|  |  |
|--|--|
|  | 0 = Reserved<br>1 = Read<br>2 = Write  |
|  | <p>Shall exist for every instruction doing memory read/write</p> <p>The record shall not be used to indicate memory accesses related to instruction address (i.e. fetch). Because we already have explicit information on instruction address (see FAQ item-3) and its instruction encoding content</p> <p>When VA→PA translation is available, this record shall follow a record with PA (not supported by the spec at this time)</p> |

| <b>InstMemContentRecord</b><br>STF_INST_MEM_CONTENT (61) |   |
|--|---|
| YES  | b[63:0] Data/content of a memory being accessed by current instruction  |
|  | <p>This happens only for memory read/write instructions</p> <p>Shall follow STF_INST_MEM_ACCESS</p> <p>For memory accesses with data size less than 8 bytes, data shall be right justified, while exact address and size are expressed in the STF_INST_MEM_ACCESS record.</p> <p>For memory accesses with data size greater than 8 bytes, use multiple STF_INST_MEM_CONTENT records, where the first STF_INST_MEM_CONTENT record refers to the address specified in the STF_INST_MEM_ACCESS record, and subsequent STF_INST_MEM_CONTENT records refer to subsequent target addresses appropriately incremented to preserve continuity of data..</p> |

| <b>BusMasterAccessRecord</b><br>STF_BUS_MASTER_ACCESS (62)                     |   |
|--|---|
| Whenever there are masters other than a single core in the traced environment. | b[63:0] Virtual address of the target memory being accessed<br>b[79:64] memory access data size<br>b[87:80] memory access initiator type<br>0 = Core<br>1 = GPU<br>2 = DMA<br>3 = PCIe<br>4 = SRIO<br>5 = ICN<br>6 = ACCEL<br>b[95:88] memory access initiator index<br>b[127:96] memory access attributes<br>b[135:128] access type<br>0 = Reserved<br>1 = Read<br>2 = Write |
|  | <p>Shall be used to capture memory accesses by masters other than the primary core being traced.</p> <p>Memory access initiator index distinguishes between multiple instances of one type of master (another core, second accelerator port, etc.)</p>  |

|  |   |
|--|---|
|  | Use Cases:<br>Another core's snoop<br>I/O device read/write |
|--|---|

| <b>BusMasterContentRecord</b><br>STF_BUS_MASTER_CONTENT (63) |   |
|--|---|
| YES  | b[63:0] Data/content of memory being accessed<br><br>This happens only for bus master read/write<br>Shall follow STF_BUS_MASTER_ACCESS<br>FAQs<br>How is information of memory access data size for STF_BUS_MASTER_ACCESS record conveyed so that we know how much valid data data is there in STF_BUS_MASTER_CONTENT?<br>Using STF_BUS_MASTER_ACCESS record's "mem access data size" field<br>How are memory access content conveyed for access data size > 8 bytes?<br>Using multiple STF_BUS_MASTER_CONTENT records. |

| <b>EventRecord</b><br>STF_EVENT (100)                                   |   |
|---|---|
| Whenever COF condition exists due to an event external to the hart/core | b[30:0] STF Exception Event ID number.<br>b[31]: Type<br>0 = Fault<br>1 = Interrupt<br>b[39:32] Number of metadata fields<br>b[103:40] Event metadata field 0<br>...<br>b[64*n+103:64*n+40] Event metadata field n<br><br>See table in the STF Exception Event and Event Meta-Data section for more details.<br><br>Shall be used to capture any type of exception or external interrupt being generated in the course of program execution. (synchronous or asynchronous).<br>EventRecord with STF_EVENT_MODE_CHANGE Event ID is a special meta-event, and shall be emitted any time there is a change in privilege mode of execution (e.g. User, Supervisor). It shall be emitted at the start of each trace to indicate the starting privilege mode and with any FORCE_PC_RECORD if there is a change in privilege mode of execution.<br><br>This record will be associated with the instruction at which the event occurs. This instruction will appear again once it successfully executes, at a later time (e.g. upon returning from an interrupt/exception handler).<br>If this event is a fault, thrown due to an invalid instruction encoding ( i.e. fetch issue), generator tool shall report it as a no-op, while preserving and reporting the PC that triggered the fault.<br><br>Event metadata fields shall be 0-extended to 64-bit<br><br>stf_dump prints this record prefixed by "EVT " |

## EventPCTargetRecord

STF\_EVENT\_PC\_TARGET (101)

|                                   |  |
|-----------------------------------|--|
| Whenever needed with an STF_EVENT | b[63:0] Virtual address of target PC when an event exception is taken, causing a change-of-flow (COF). |
|                                   | May follow STF_EVENT<br>Provides information on which PC should follow this event.                     |

## InstMicroOpRecord

STF\_INST\_MICROOP (230)

|  |  |
|--|--|
| Whenever STF_CONT AIN_MICR OOP bit is set in STF_TRACE_INFO record | b[7:0] Size of micro-op<br>b[39:8] Micro-op  |
|  | Use-case:<br>Used by trace morphing tools to inject new instruction encoding records, a technique used to run micro-architecture "what-if" studies when original instruction is replaced by more than one instruction<br>UseCase 1: Instruction replacement 1:1 in which case trace's replaced instruction would be indicated by STF_INST_32 record. In this case there is no microop record emitted in the trace<br>UseCase 2: Instruction replacement 1:n where n = (#of microops emitted + 1).<br>In both cases following hold true:<br>The original Instruction record STF_INST_MICROOP replacement and/or microop related changes to memory accesses, operand reg records etc. are emitted(or removed) in the trace before the original STF_INST_32/16 record that defines instruction boundary for all artifacts related to the instruction(IRG) |

## Inst32Record

STF\_INST\_32 (240)

|  |   |
|--|---|
| Whenever instruction is a 32-bit instruction | b[31:0] 32-bit instruction encoding   |
|  | Shall conform to the last STF_INST_IEM record before this record<br>All preceding instruction attributes are pulled together with this record to create an instruction record group (IRG (i.e. STFInst)). |

## Inst16Record

STF\_INST\_16 (241)

|  |   |
|--|---|
| Whenever instruction is a 16-bit instruction | b[15:0] 16-bit instruction encoding   |
|  | Shall conform to the last STF_INST_IEM record before this record.<br>All preceding instruction attributes are pulled together with this record to create an instruction record group (IRG (i.e. STFInst)) |

|   |   |
|---|---|
| <b>STFReserveEndRecord</b><br>STF_RESERVE_END (255) |   |
|   | b[-1] Record has no data, just the descriptor.                          |
|   | This shall be the last record of a trace. Reserved for error detection. |

# STF Exception Event and Event Meta-Data

The following shows definitions of various STF exception event identifiers and associated meta-data for the events captured by an EventRecord (STF\_EVENT).

| STF Exception ID              | Exception Related Data/Content shall Include   | Sync/Async Exception | Description                         |
|-------------------------------|--|----------------------|-------------------------------------|
| STF_EVENT_USER_SOFTWARE       | Field 0: Source of the interrupt   | Async                | User software interrupt             |
| STF_EVENT_SUPERVISOR_SOFTWARE |  | Async                | Supervisor software interrupt       |
| STF_EVENT_HYPERVISOR_SOFTWARE |  | Async                | Hypervisor software interrupt       |
| STF_EVENT_MACHINE_SOFTWARE    |  | Async                | Machine software interrupt          |
| STF_EVENT_USER_TIMER          |  | Async                | User timer interrupt                |
| STF_EVENT_SUPERVISOR_TIMER    |  | Async                | Supervisor timer interrupt          |
| STF_EVENT_HYPERVISOR_TIMER    |  | Async                | Hypervisor timer interrupt          |
| STF_EVENT_MACHINE_TIMER       |  | Async                | Machine timer interrupt             |
| STF_EVENT_USER_EXT            |  | Async                | User external interrupt             |
| STF_EVENT_SUPERVISOR_EXT      |  | Async                | Supervisor external interrupt       |
| STF_EVENT_HYPERVISOR_EXT      |  | Async                | Hypervisor external interrupt       |
| STF_EVENT_MACHINE_EXT         |  | Async                | Machine external interrupt          |
| STF_EVENT_COPROCESSOR         |  | Async                | Supervisor guest external interrupt |
| STF_EVENT_INST_ADDR_MISALIGN  | Field 0: Virtual address of the instruction  | Sync                 | Instruction Address Misaligned      |
| STF_EVENT_INST_ADDR_FAULT     | Field 0: Virtual address of the instruction  | Sync                 | Instruction access fault            |
| STF_EVENT_ILLEGAL_INST        | Field 0: Virtual address of the instruction<br>Field 1: Instruction opcode<br>Field 2: Instruction encoding mode | Sync                 | Illegal instruction                 |
| STF_EVENT_BREAKPOINT          | Field 0: Virtual address of the instruction  | Sync                 | Breakpoint                          |
| STF_EVENT_LOAD_ADDR_MISALIGN  | Field 0: Virtual address of the instruction  | Sync                 | Load address misaligned             |

|                               |  |            |   |
|-------------------------------|--|------------|---|
| STF_EVENT_LOAD_ACCESS_FAULT   | Field 1: Instruction encoding  | Sync       | Load access fault   |
| STF_EVENT_STORE_ADDR_MISALIGN | Field 2: Target address that caused the exception  | Sync       | Store address misaligned                                      |
| STF_EVENT_STORE_ACCESS_FAULT  |  | Sync       | Store access fault  |
| STF_EVENT_USER_ECALL          | Field 0: System call number  | Sync       | Environment call from User mode                               |
| STF_EVENT_SUPERVISOR_ECALL    |  | Sync       | Environment call from Supervisor mode                         |
| STF_EVENT_HYPERVISOR_ECALL    |  | Sync       | Environment call from Hypervisor mode                         |
| STF_EVENT_MACHINE_ECALL       |  | Sync       | Environment call from Machine mode                            |
| STF_EVENT_INST_PAGE_FAULT     | Field 0: Virtual address of the instruction<br>Field 1: Instruction encoding   | Sync       | Instruction page fault  |
| STF_EVENT_LOAD_PAGE_FAULT     | Field 0: Virtual address of the instruction  | Sync       | Load page fault   |
| STF_EVENT_STORE_PAGE_FAULT    | Field 1: Instruction encoding<br>Field 2: Target address that caused the exception   | Sync       | Store page fault  |
| STF_EVENT_MODE_CHANGE         | Field 0: Mode in which subsequent instructions execute<br>Field 0 Encoding:<br>00 - User Mode<br>01 - Supervisor Mode<br>10 - Hypervisor Mode<br>11 - Machine Mode | Sync/Async | Meta-event for capturing transition into a new execution mode |

# Terminology / Definitions

**Change of Flow (COF)** - Change of instruction execution flow for reasons that are deterministic as well as non-deterministic. COF is logged in the trace using the STF\_INST\_PC\_TARGET record for deterministic and with STF\_FORCE\_PC , STF\_INST\_EXCEPTION\_PC\_TARGET and STF\_EVENT\_PC\_TARGET for non-deterministic cases.

STF\_FORCE\_PC applies COF to the current instruction, while all other cases apply it to the subsequent instruction.

1. Deterministic COF cases: These COF cases happen within the scope of a thread context and without having any exceptions being raised within that context.
  - a. Function returns through "ret" like instruction
  - b. Direct branches through branch instructions where target address of where the execution flow will jump to an address which is either a PC relative(relative to current address) or absolute address
  - c. Indirect branches through branch instructions where the target address of where the execution flow will jump to is an address stored in a register. The address was placed in that register by calculation by earlier instruction/s. Typical use case of this is jump to a function called by a function pointer.
2. Non-deterministic COF cases: These COF cases happen due to interruptions in a program thread due to exceptions being raised either by the program thread context itself or due to external interruptions. This category of COF is indicated by
  - a. Program flow changes, due to actions of current instruction, to exception handler within the same exception level. Example can be kernel code accessing data from address whose entry is missing in PTE, thus triggering page fault exception within the same exception level.
  - b. Program flow changes, due to actions of current instruction, to exception handler to a higher exception level. Example: user-space program trying to access memory causing page fault causing exception level changes to supervisor mode or user-space syscall.
  - c. Program flow changes due to return from exception handler
  - d. Program flow changes, due to context switches. Example: thread switch, process switch, Address Space Identifier (ASID) switch.
  - e. Program flow changes, due to external interrupts to the same or higher exception level.
  - f. Program flow changes due to instruction mode change from this set (not applicable to RISC-V)
  - g. Program flow changes due to manipulation of a trace file.